Does the flatter runtime curve hold?

Measured end-to-end cost through 65,536 orbit queries · lower is better

CONTROLLED STACK · 1,024 simulated 33×33 exposures over 8 days

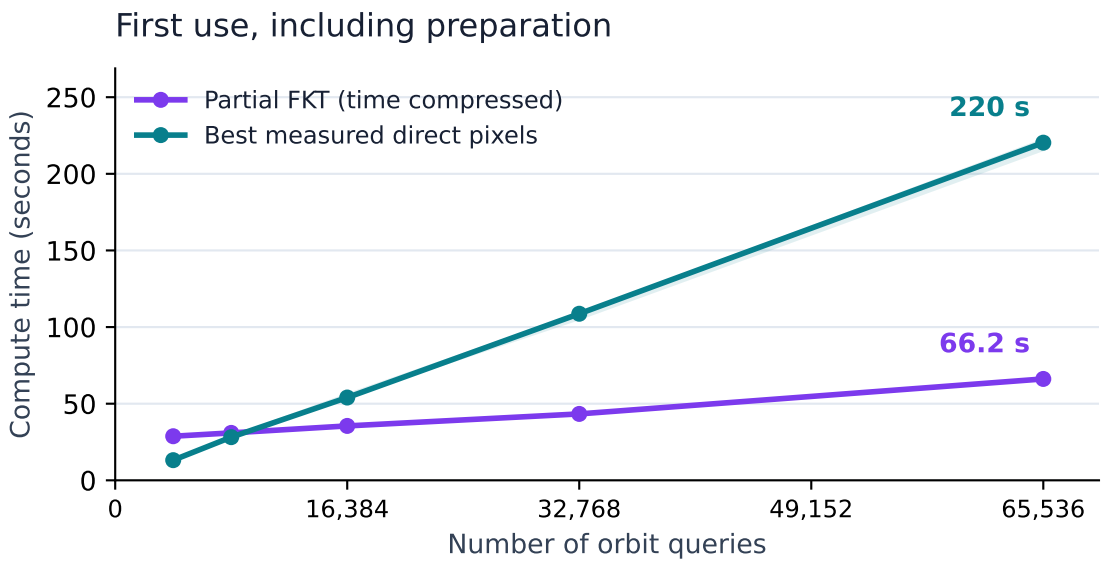

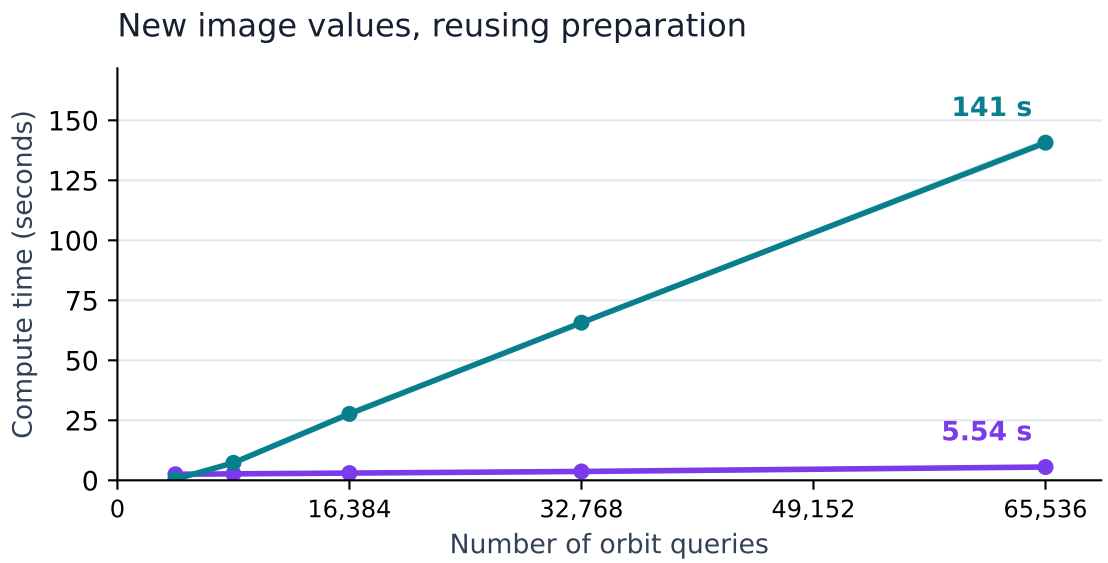

REAL PIXELS · 3 ZTF exposures over 12.3 minutes

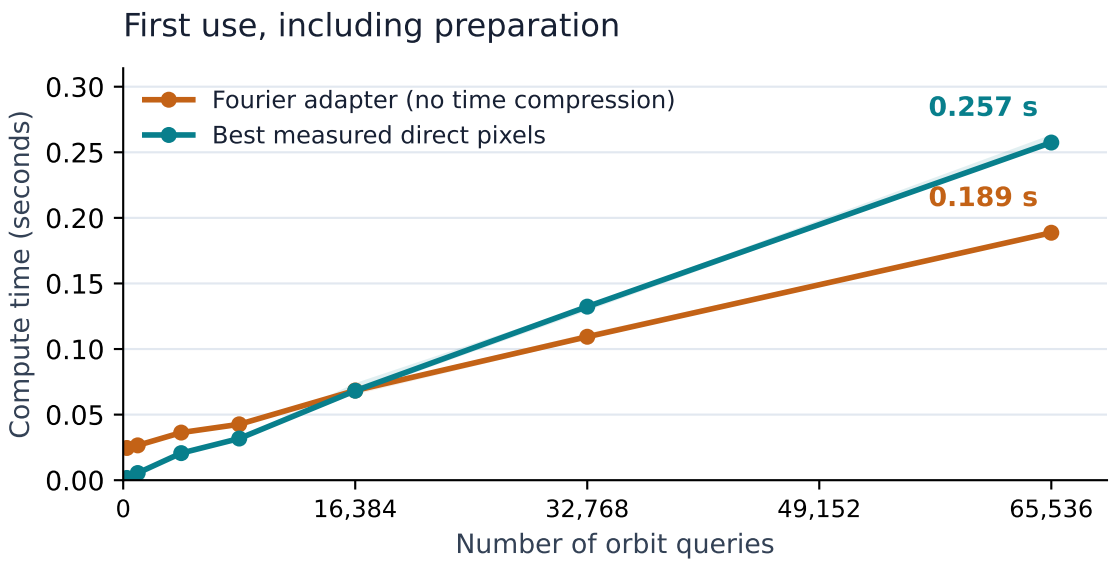

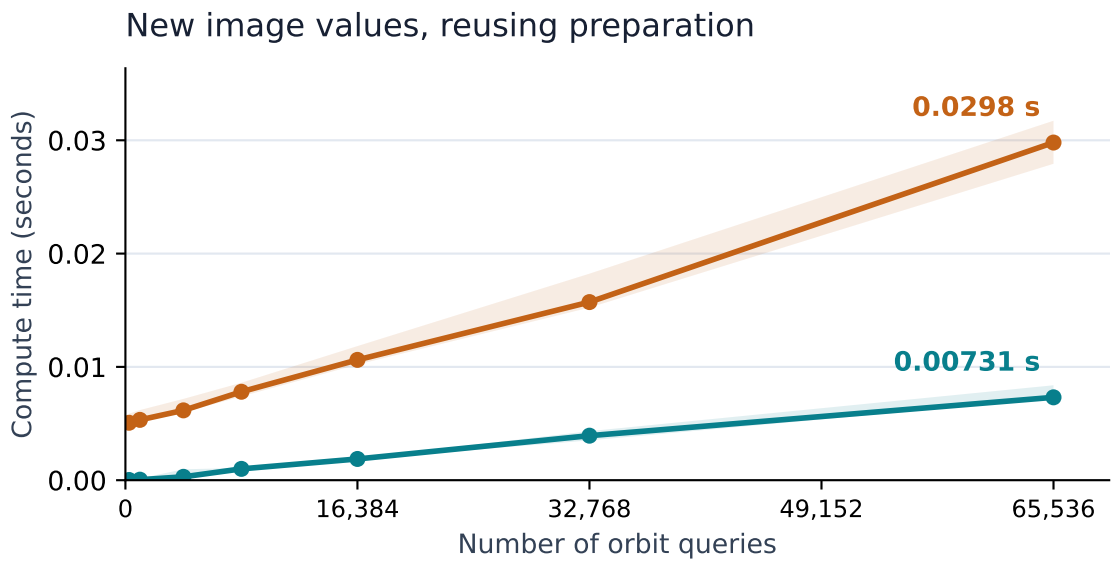

Points use measured medians; shading spans application samples (not confidence intervals). Setup measured once. Panel scales differ.
Direct methods retain up to 3 GB of cache, recomputing uncached PSF factors. Both routes reuse information on new images.
The real-data row tests the Fourier image adapter only. It does not validate time compression on real ephemerides.